

Characterizing the object categories in two children's home visual environment across development

Anonymous CogSci submission

Abstract

What do infants and young children tend to see in their everyday lives? Relatively little work has examined the categories and objects that tend to be in the infant view during everyday experience, despite the fact that this knowledge is central to theories of category learning. Here, we analyzed the prevalence of the superordinate categories (e.g., people, animals, food) in the infant view in a longitudinal dataset of egocentric infant visual experience. Overall, we found a surprising amount of consistency in the broad characteristics of children's visual environment across individuals and across developmental time, in contrast to prior work examining the changing nature of the social signals in the infant view. In addition, we analyzed the distribution and identity of the categories that children tended to touch and interact with in this dataset, generalizing previous findings that these objects tended to be distributed in a Zipfian manner. Taken together, these findings take a first step towards characterizing infants' changing visual environment, and call for future work to examine the generalizability of these results and to link them to learning outcomes.

Keywords: Object categorization; infant visual experience; head-mounted cameras; longitudinal data.

Introduction

What do children tend to see in their everyday lives? While an understanding of children's visual environment is central to both theories of language acquisition and visual development, we know remarkably little about the categories and objects that tend to be in the infant view, or in what format they are experienced. For example, how often do infants tend to see animals in real-life vs. in storybooks or as toys? How consistent are children's visual environments across individuals and across developmental time?

Over the past decade, researchers have begun to answer these questions by documenting the infant egocentric perspective using head-mounted cameras (Franchak, Kretch, Soska, & Adolph, 2011; Yoshida & Smith, 2008) and quantifying the degree to which there are substantial shifts in infants' viewpoints that may have downstream developmental consequences. As adults, it is hard to intuit how strange this viewpoint can be, and how much it varies across development, transitioning over the first two years of life from close-up views of faces to restricted views of hands manipulating objects (Fausey, Jayaraman, & Smith, 2016; Long, Kachergis, Agrawal, & Frank, 2020), with children's postural developments to a large extent shaping what they see (Sanchez, Long, Kraus, & Frank, 2018). Most work, however, has focused on documenting the social information that infants and

children have access to across early development (Fausey et al., 2016; Sanchez et al., 2018; Yoshida & Smith, 2008).

More recent research has made progress towards understanding what objects tend to be in the infant view, starting with analyzing the basic-level categories (e.g., spoons, cups) in the view of 8-month-olds during mealtime. This work suggests that a small number of objects are both pervasively present during mealtime and among infants' first-learned words (Clerkin, Hart, Rehg, Yu, & Smith, 2017), pointing towards a link between visual experience and early word learning and category learning.

Thus, a more complete understanding of the visual environment of infants and young children could yield insights about the inputs to both category learning and word learning. Indeed, different distributions of these visual referents lead to constraints on the kinds of learning mechanisms that must operate to form robust category representations – and to learn words for these categories. However, at present, no datasets are sufficiently annotated to constrain these theoretical accounts.

For example, if the categories in the infant view shift dramatically over the first few years of life, then we might expect infants to learn about certain categories earlier vs. later during development. Prior work documenting the proportion of social information in view has suggested that children see more hands relative to faces in this same age range (Fausey et al., 2016; Long et al., 2020). Thus, one possibility is that as children learn to crawl and walk (Franchak et al., 2011; Long et al., 2020; Sanchez et al., 2018), categories that children are likely to interact with (i.e., toys, small objects) may also become more prevalent in the child's view. If this was the case, this finding would support a view where the inputs to early category learning are shaped by children's own ability to actively explore their environment.

On the other hand, the broad characteristics of children's visual environments may be relatively stable and mostly determined by the activities that they tend to engage in. Indeed, some theoretical accounts have suggested that the statistics of children's visual environment are mostly driven by these stereotyped activity contexts (Bruner, 1985) – e.g. mealtime or storytime – and that children learn most robustly in these contexts. On these accounts, children might become very sensitive to the co-occurrences between different activities (e.g., eating) and object categories (e.g., spoons, food). However,



Figure 1: Example frames with annotations of four different broad categories.

no work has identified how consistently categories co-occur in natural environments. For example, while some activity contexts (e.g., storytime) lead to intuitive co-occurrences between object categories (e.g., between books and people), not all activity contexts will generate intuitive or consistent co-occurrences between object categories.

Finally, *how* infants interact with object categories will undoubtedly change what they learn about them. For example, children tend to generate informative views of objects while manipulating them – and, early in development, children’s ability to sit and manipulate objects correlates with their perceptual abilities (Soska, Adolph, & Johnson, 2010). Yet while most datasets used to train deep neural network models contain only photographs of object categories, many children – especially those in Western, industrialized cultures – will likely experience many categories through picture books as flat, stylized, 2D depictions. If children see very few real-life exemplars of a category relative to depictions (e.g., giraffes), this suggests that children must learn to generalize between these different visual formats in order to group these exemplars into one category – and, further, implies that these representations might be coarser than those experienced across many different formats. And if children only interact and manipulate a few small set of categories – as suggested by Clerkin et al., 2017 – children may first learn about these frequently experienced categories and then use these representations to generalize to the categories they encounter very infrequently.

Here, we take a step towards answering these questions by characterizing the visual environment of two young children in a longitudinal corpus of head-mounted camera data (Sullivan, Mei, Perfors, Wojcik, & Frank, 2020) from 6-32 months of age. To characterize broad trends in the visual environment over development, we analyzed the superordinate categories

of objects (e.g., animals, vehicles, toys, people) present in the infant view, obtaining annotations on a randomly sampled set of 24,000 frames. To provide a closer look into the kinds of objects children have the most intensive visual and haptic experience with, we also examined the basic-level categories that children interacted with during these everyday activities. To do so, we annotated the basic-level categories of the objects that children were interacting with in the subset of frames where children’s hands were visible.

Method

Dataset

The dataset is described in detail in Sullivan et al. (2020). Children wore Veho Muvi miniature cameras mounted on a custom camping headlamp harness (“headcams”) at least twice weekly, for approximately one hour per recording session. One weekly session was on the same day each week at a roughly constant time of day, while the other(s) were chosen arbitrarily at the participating family’s discretion. At the time of the recording, all three children were in single-child households. Videos captured by the headcam were 640x480 pixels, and a fisheye lens was attached to the camera to increase the field of view to approximately 109 degrees horizontal x 70 degrees vertical. We randomly sampled 24,000 frames from videos of two of the children in the dataset (S, A) over the entire age range.

Annotation procedures

Broad categories in view Annotations of the broad categories in the dataset were obtained using AWS Sagemaker annotations. Participants were instructed to select all of the categories that could be applied to an image; two workers annotated each image, and each category that was annotated for an image was assigned a confidence score (possible range:

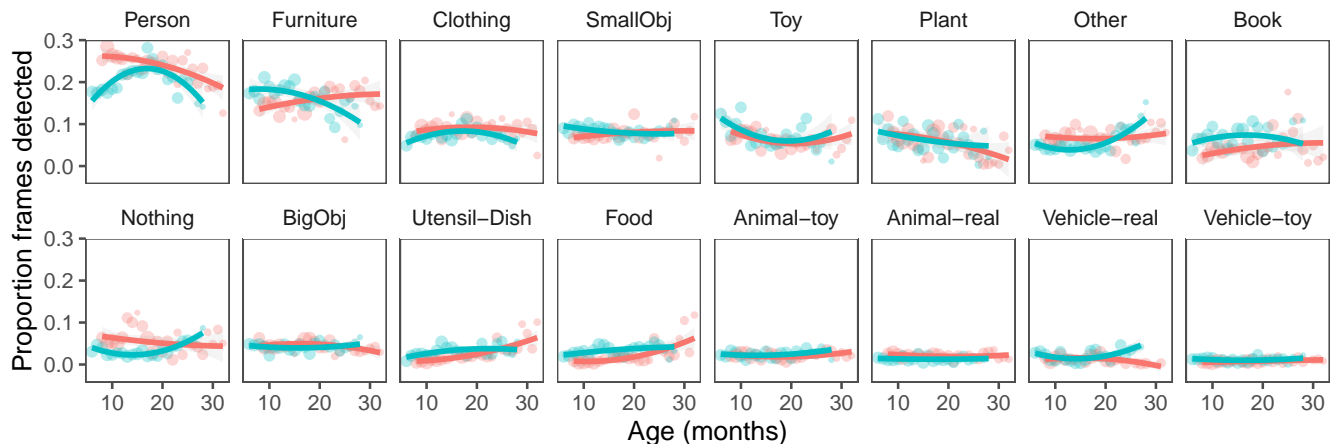


Figure 2: Frequency of categories annotated across the 24K random frames plotted as a function of each child’s age (in months); each child’s age was calculated in days relative to the date that the videos were filmed and converted to months. Each color represents data from a different child.

0-1, range in dataset: .5-1). Participants selected whether the following categories were present in the shown image: Animal (real), Animal (toy/drawing), Vehicle (real), Vehicle (toy/drawing), Plant, Clothing, Person, Furniture, Food, Utensil/Dish, Other Small Object, Other Big Object, Book, Other, or Nothing visible. We included “other small objects” and “other big objects” as categories that participants could use to indicate objects that fell outside of these traditional superordinate categories (i.e. furniture, plant, toy). Additional instructions were provided to specify that ‘other big object’ refers to an object that is bigger than a chair, and that ‘other small object’ refers to an object that is small enough to be held with one or two hands (Konkle & Oliva, 2012). Annotators were required to select at least one category before proceeding. Individual annotations had confidence scores below the 25th percentile were excluded from analyses (although all conclusions hold with and without these low-confidence annotations).

We assessed reliability of these annotations by comparing them to annotations made on the same task for a random subsample of 1200 frames on AWS Sagemaker, again using two workers per image. When excluding low-confidence annotations (for a final sample of $N=950$ images), we found agreement was moderate (average Cohen’s Kappa = 0.29, but varied substantially between different categories (range = 0.02, 0.59), with an average disagreement rate of 11.41% across categories. Annotators disagreed most on whether “clothing” was present in an image and whether “other big object” was present (i.e. a big object that was not a piece of furniture or a vehicle). To assess the nature of these disagreements, we manually examined a random sample of 160 images with disagreements with 10 images from each category. There will be relatively equal proportions of images where annotations failed to identify a clear example of a category ($M=25\%$) or where annotations select an erroneous category label ($M=24\%$). However, we found that most ($M=50.62\%$)

of the disagreements resulted from ambiguous exemplars, for example where the category was present but very distant, occluded, or blurry. Annotators also showed some disagreement about whether glossy photos of different categories in books should be counted as “real” or “toy/drawing,” and whether partial views of people (i.e. child’s own hands) should count as a “person.” Going forward, we analyze the larger set of annotations with the caveat that there is inevitably some ambiguity in what counts as an exemplar of these categories and that these data include both misses and false alarms. All annotations and images are openly available on the repository for this associated project ([OSF LINK HERE.](#))

Objects in the child’s reach We also annotated the objects that children were interacting with in a subset of these frames. First, we selected the frames in which participants recruited via Amazon Mechanical Turk indicated that a child’s hand or hands were visible in the image (see Long et al., 2020) and one author annotated 1817 of these frames, spanning 7 to 28 months of age with roughly equal proportions from the two children. The annotator noted what object the child was touching or pointing to with in frames containing children’s hands, using basic level object categories such as “block” and “cracker.” If a child was holding a book and pointing to a depicted object in the book, the depicted object was noted as the category they were interacting with; otherwise, it was noted as ‘book.’ When children were interacting with drawing or toy versions, these annotations were marked with a ‘-drawing’ and ‘-toy’ modifier and counted as separate entries. If a view was allocentric, there were no child hands in view, or there were no objects that were visible, these frames were excluded from analysis; this left 1313 frames with annotations.

Results

Which categories are prevalent in the child's view?

First, we examined the overall prevalence of each broad category in the infant view. Somewhat surprisingly, we found that the prevalence of most of these categories were relatively stable both across the two children in the dataset as well as over developmental time. This stands in contrast to prior work on the prevalence of faces/hands in the infant view (Fausey et al., 2016; Long et al., 2020), suggesting that these broader characteristics of children's visual experience may be more consistent.

We next examined the details of these environments. We found that people were by far the most prevalent of these categories: over 20% of the annotated frames contained people, far more than any other category (including all kinds of toys combined). In contrast, there were relatively few instances of animals in the infant view—either as toys or their real-life counterparts. Less than 5% of the frames contained any kind of depicted or real animal, and those few frames that did contained depicted vs. real animals in equal proportion. Manual inspection of these frames containing animals revealed that the “real” animals had relatively little variety – they were overwhelmingly frames containing images of household pets (i.e., cats, dogs, and chickens, in the case of A), whereas the animals that were “toys/drawings” depicted a much larger variety of animals, as one might expect. Overall, these results suggest that – at least for these children – people are much more frequent than depictions or real-life versions of animals, indicating that toys and drawings may provide frequent input to their representations of these categories – despite the fact that animal names are often among children's first words (Frank, Braginsky, Yurovsky, & Marchman, 2020) and often referenced in storybooks.

Far more prevalent than animals, instead, were objects. Views of furniture were the next most common category after “people”. However, in older age ranges, “big objects” – including furniture, vehicles, and other big objects – tended to be less frequently in the view of infants than “small” objects – including toys (of all kinds), food, utensils, books, and other small objects (see Figure 3). This effect was much exaggerated when we conducted this analysis on a subset of the frames where children's hands were also in view as a proxy for times when children were interacting with objects. In these frames, small objects tended to be much more prevalent than the frames that we annotated. These data are consistent with the idea that as children grow and become more adept at handling objects on their own, small objects may tend to be more often in view.

Which categories co-occur in children's visual environment?

Next, we next examined the degree to which these broad categories appeared together in different frames. Figure 4 shows the co-occurrence of the broad categories, and reveals some relatively intuitive patterns that may reflect activity con-

texts. For example, “dishes” and “food” co-occurred quite frequently together, as did “people” and “clothing,” and most animals that were toys or drawings appeared when “books” were also present. To determine which cells significantly deviate from chance we used a permutation analysis in which we shuffled the annotated category labels within each frame and examined the distribution of co-occurrences across 100 randomized co-occurrence matrices. The cells in the plot that occurred fewer times than expected by chance (<99% of permuted cells) are labeled with a ‘-’, while those that occurred more often than expected by chance (>99% of permuted cells) are labeled with a ‘+’. Broadly, these results suggest that activity contexts – such as playtime, mealtime, or storytime – may have a considerable influence on structuring the categories that tend to be in the infant view, pointing towards the role of these contexts for shaping what children learn about these categories (Bruner, 1985).

What objects do children tend to interact with?

While many different categories may be in the child's view, not all of these objects may be experienced in the same way. In particular, it may be that children are more likely to form robust representations of objects that they physically interact with more often, and by extension they may also learn the labels of these objects earlier. In this analysis, we sought to analyze the basic-level identities of the objects that children tended to be interacting with in their home environments, and the distributions of those identities. While some work has found that the objects in view during mealtime tend to have a Zipfian distribution (Clerkin et al., 2017), it is not yet known whether this finding will extend to objects that do not appear during mealtime and that children interact with during a wide range of activities. For example, there may be far fewer objects that are only interacted with a limited number of times vs. seen a limited number of times.

First, we found that the distribution of the objects in view roughly followed a power law distribution (with $\alpha = 1.82$), confirming that the distribution of the objects that children interact with in general (not only during mealtime as measured by Clerkin et al., 2017) is highly skewed – as is the distribution of categories that they tend to see. In the 148 frames with objects that children were interacting with, we found 132 unique categories when collapsing across formats (i.e. drawings, toys, real-life), and 148 when exemplars were considered separately across formats. When we examined which categories were most frequent, we found that books were overwhelmingly the most present object in the views of these two children, comprising over 15% of the objects that children were seen to be manipulating. Generic baby toys (that were unidentifiable to the authors as specific toys) were the next most frequent object category, and children were often seen to be touching or holding on to their caregivers (see top 20 most frequent categories in Figure 5). Additionally, frames in which objects were occluded or unidentifiable had a high frequency of appearance, accounting for around 6% of frames, affirming previous assertions that infant egocentric

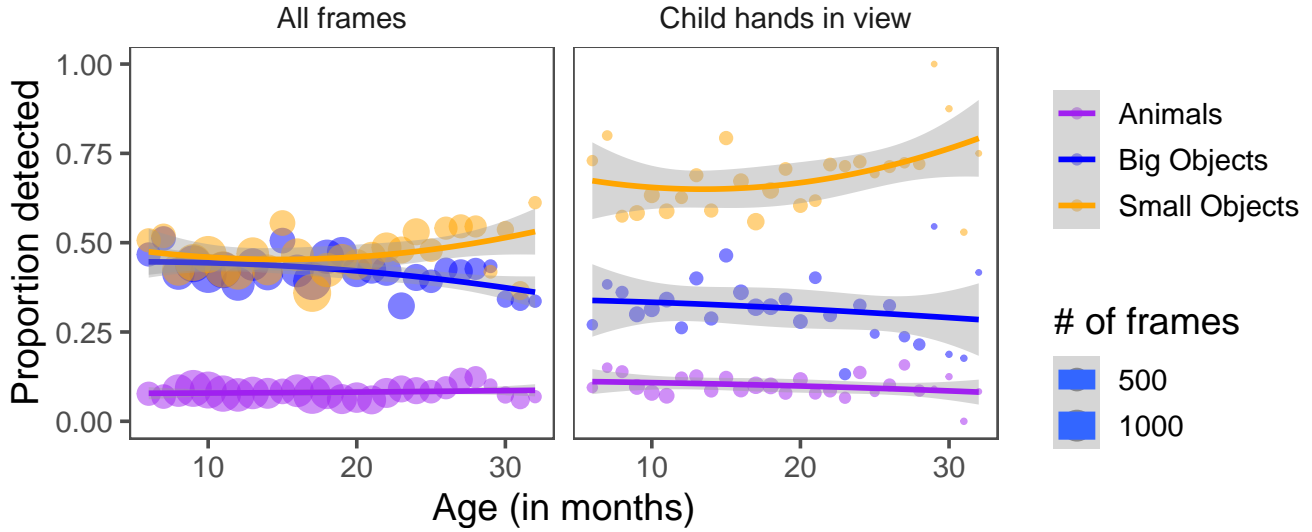


Figure 3: Frequency of animals (including toys) relative to big and small inanimate objects detected in the dataset, both when analyzing all frames that were annotated (left) and the subset of frames where a child’s hand was visible in the frame (right).

views are nonintuitive.

General Discussion

What determines the categories that infants tend to see and interact with across early development? To examine the categories in the infant view, we analyzed a sample of random frames taken from a longitudinal dataset of two children (Sullivan et al., 2020). Overall, we found relative consistency in children’s visual environment over development, in contrast to prior work on the prevalence of social signals over this same developmental time period (Fausey et al., 2016; Long et al., 2020). The relative proportions of broad categories of objects (i.e., furniture, toys, animals, people) was relatively consistent among the two individuals here, and across developmental time. Instead, this work suggests that activity contexts, such as mealtime or storytime (Bruner, 1985) may structure the broad characteristics of young children’s visual environment and the co-occurrences between different object categories; data-driven analysis of these category co-occurrence revealed stereotypical combinations (i.e. utensils and food, people and clothing).

However, while people were incredibly frequent in the child’s view, animals – either as toys or their real-life versions – were relatively infrequent and occurred in equal proportions. This stands in contrast to a long literature documenting that even newborns prefer to attend to animate agents (Farroni et al., 2005), that visual cortex dedicates a remarkable amount of space to processing animals (Konkle & Caramazza, 2013), and that animal names tend to be among children’s first-learned words (Frank et al., 2020). Therefore, these results underscore that children’s heightened attention to animals (Farroni et al., 2005) likely interacts with frequency of occurrence in the visual field to drive early category learning.

Instead, we found the child’s view was most likely to be

dominated by small objects that they were interacting with – such as food, books, or toys. We also found that the proportion of these small objects increased dramatically when we restricted our analysis to frames where the child’s hands are in view, suggesting that the statistics of children’s visual environment shift substantially when they are acting on the world themselves (i.e., while playing). More generally, we found that the distribution of these objects seem to follow a Zipfian distribution, as does word usage in natural language. Indeed, as mealtime has previously been used to characterize the objects in the infant view (Clerkin et al., 2017) and frames with food or utensil and dishes accounted for less than 5% of views in the SAYcam dataset, we were unsure whether this would be the case. However, this analysis suggests that – at least in this sample – that infants may more generally interact with different object categories in a relatively Zipfian manner.

Overall, this work takes a first step in characterizing the categories in the visual environment of children over development, calling for future work to understand the generalizability of these findings beyond the present dataset. While we found relatively consistent results across age and the two children in the dataset, both of these children are from relatively similar households and cultural contexts. Nonetheless, we predict broad generality of the findings that objects are more frequent than animals and that categories will be relatively stable in their prevalence across age. In particular, we predict that most children in urban or suburban contexts are unlikely to see real animals more frequently than depicted animals, and the distribution of objects that children interact with are likely to follow a Zipfian distribution – regardless of which specific objects these are.

More broadly, this work highlights the need for systematic investigations of how the frequency of the categories in the child’s view interacts with different attentional biases, learn-




Figure 4: Co-occurrence between different categories detected in the dataset across all frames. Each cell represents the probability that the category on the y-axis (e.g., clothing) occurs relative to the occurrence of the category on the x-axis (e.g., person). Lighter values indicate higher probabilities of co-occurrence (max=.8, min=0). A permutation analysis was used to determine which cell values were outside the 99% confidence interval of counts: - = less than 99%, + = greater than 99%.

ing mechanisms, and social cues to produce robust representations that support early category and language learning. An understanding of what is – and what is not – learnable solely from frequent exposures will provide constraints on our accounts of the learning mechanisms that allow children to learn so much so quickly.

Acknowledgements

(Blinded)

References

Bruner, J. (1985). The role of interaction formats in language acquisition. In *Language and social situations* (pp. 31–46). Springer.

Clerkin, E. M., Hart, E., Rehg, J. M., Yu, C., & Smith, L. B. (2017). Real-world visual statistics and infants' first-learned object names. *Phil. Trans. R. Soc. B*, 372(1711), 20160055.

Farroni, T., Johnson, M. H., Menon, E., Zulian, L., Faraguna, D., & Csibra, G. (2005). Newborns' preference for face-relevant stimuli: Effects of contrast polarity. *Proceedings of the National Academy of Sciences*, 102(47), 17245–17250.




Figure 5: Top 20 most frequent categories that children's hands were interacting with in these egocentric videos.

Fausey, C. M., Jayaraman, S., & Smith, L. B. (2016). From faces to hands: Changing visual input in the first two years. *Cognition*, 152, 101–107.

Franchak, J. M., Kretch, K. S., Soska, K. C., & Adolph, K. E. (2011). Head-mounted eye tracking: A new method to describe infant looking. *Child Development*, 82(6), 1738–1750.

Frank, M. C., Braginsky, M., Yurovsky, D., & Marchman, V. A. (2020). *Variability and Consistency in Early Language Learning: The Wordbank Project*. Cambridge, MA: MIT Press. Retrieved from <http://wordbank-book.stanford.edu>

Konkle, T., & Caramazza, A. (2013). Tripartite organization of the ventral stream by animacy and object size. *Journal of Neuroscience*, 33(25), 10235–10242.

Konkle, T., & Oliva, A. (2012). A familiar-size stroop effect: Real-world size is an automatic property of object representation. *Journal of Experimental Psychology: Human Perception and Performance*, 38(3), 561.

Long, B., Kachergis, G., Agrawal, K., & Frank, M. C. (2020). Detecting social information in a dense database of infants' natural visual experience. <https://Psyarxiv.com/Z7tdg/>.

Sanchez, A., Long, B., Kraus, A. M., & Frank, M. C. (2018). Postural developments modulate children's visual access to social information. In *Proceedings of the 40th annual conference of the cognitive science society*.

Soska, K. C., Adolph, K. E., & Johnson, S. P. (2010). Systems in development: Motor skill acquisition facilitates three-dimensional object completion. *Developmental Psychology*, 46(1), 129.

Sullivan, J., Mei, M., Perfors, A., Wojcik, E., & Frank, M. C. (2020). SAYCam: A large, longitudinal audiovisual dataset recorded from the infants perspective. *PsyArXiv*. Retrieved from <https://psyarxiv.com/fy8zx/>

Yoshida, H., & Smith, L. B. (2008). What's in view for toddlers? Using a head camera to study visual experience. *Infancy*, 13(3), 229–248.